

assignment12_MeyerJake

May 29, 2023

0.1 Assignment 12

0.1.1 DSC 650

0.1.2 Jake Meyer

0.1.3 05/29/2023

Using code examples from Chapter 6 of First Edition: [deep-learning-with-python-notebooks](#)

```
[2]: ## Import the necessary modules for the assignment.
import csv
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from keras.datasets import mnist
import tensorflow as tf
import keras
import sklearn
import itertools
import string
from numpy import array
from numpy import argmax
from sklearn.model_selection import train_test_split
from scipy.stats import norm
from pathlib import Path
from contextlib import redirect_stdout
import time
import os
import random
import sys

## Import the necessary keras components.
from keras import layers, models, preprocessing
from keras.datasets import imdb
from keras.preprocessing.text import Tokenizer
from keras.preprocessing import sequence
from keras.utils import to_categorical, np_utils
from keras.models import Sequential, load_model, Model
from keras.layers import LSTM
```

```

from keras.layers.core import Dense, Dropout, Activation, Flatten, Embedding
from keras.optimizers import RMSprop
from keras import backend as K
import tensorflow.compat.v1 as tf
tf.disable_v2_behavior()

```

```

[3]: ## Print versions of essential packages
print("keras version: {}".format(keras.__version__))
print("tensorflow version: {}".format(tf.__version__))
print("pandas version: {}".format(pd.__version__))
print("numpy version: {}".format(np.__version__))

```

```

keras version: 2.11.0
tensorflow version: 2.11.0
pandas version: 1.5.3
numpy version: 1.24.2

```

0.2 Assignment 12 Criteria

Using section 8.4 in Deep Learning with Python as a guide, implement a variational autoencoder using the MNIST data set and save a grid of 15 x 15 digits to the results/vae directory. If you would rather work on a more interesting dataset, you can use the CelebFaces Attributes Dataset instead.

```

[4]: '''
Example Code from 8.4 in Deep Learning with Python provided below for reference:
*****
# Encode the input into a mean and variance parameter
z_mean, z_log_variance = encoder(input_img)

# Draw a latent point using a small random epsilon
z = z_mean + exp(z_log_variance) * epsilon

# Then decode z back to an image
reconstructed_img = decoder(z)

# Instantiate a model
model = Model(input_img, reconstructed_img)

# Then train the model using 2 losses:
# a reconstruction loss and a regularization loss
*****
img_shape = (28, 28, 1)
batch_size = 16
latent_dim = 2 # Dimensionality of the latent space: a plane

input_img = keras.Input(shape=img_shape)

```

```

x = layers.Conv2D(32, 3,
                  padding='same', activation='relu')(input_img)
x = layers.Conv2D(64, 3,
                  padding='same', activation='relu',
                  strides=(2, 2))(x)
x = layers.Conv2D(64, 3,
                  padding='same', activation='relu')(x)
x = layers.Conv2D(64, 3,
                  padding='same', activation='relu')(x)
shape_before_flattening = K.int_shape(x)

x = layers.Flatten()(x)
x = layers.Dense(32, activation='relu')(x)

z_mean = layers.Dense(latent_dim)(x)
z_log_var = layers.Dense(latent_dim)(x)
*****
def sampling(args):
    z_mean, z_log_var = args
    epsilon = K.random_normal(shape=(K.shape(z_mean)[0], latent_dim),
                              mean=0., stddev=1.)
    return z_mean + K.exp(z_log_var) * epsilon

z = layers.Lambda(sampling)([z_mean, z_log_var])
*****
# This is the input where we will feed `z`.
decoder_input = layers.Input(K.int_shape(z)[1:])

# Upsample to the correct number of units
x = layers.Dense(np.prod(shape_before_flattening[1:]),
                  activation='relu')(decoder_input)

# Reshape into an image of the same shape as before our last `Flatten` layer
x = layers.Reshape(shape_before_flattening[1:])(x)

# We then apply then reverse operation to the initial
# stack of convolution layers: a `Conv2DTranspose` layers
# with corresponding parameters.
x = layers.Conv2DTranspose(32, 3,
                           padding='same', activation='relu',
                           strides=(2, 2))(x)
x = layers.Conv2D(1, 3,
                  padding='same', activation='sigmoid')(x)
# We end up with a feature map of the same size as the original input.

# This is our decoder model.
decoder = Model(decoder_input, x)

```

```

# We then apply it to `z` to recover the decoded `z`.
z_decoded = decoder(z)
*****
class CustomVariationalLayer(keras.layers.Layer):

    def vae_loss(self, x, z_decoded):
        x = K.flatten(x)
        z_decoded = K.flatten(z_decoded)
        xent_loss = keras.metrics.binary_crossentropy(x, z_decoded)
        kl_loss = -5e-4 * K.mean(
            1 + z_log_var - K.square(z_mean) - K.exp(z_log_var), axis=-1)
        return K.mean(xent_loss + kl_loss)

    def call(self, inputs):
        x = inputs[0]
        z_decoded = inputs[1]
        loss = self.vae_loss(x, z_decoded)
        self.add_loss(loss, inputs=inputs)
        # We don't use this output.
        return x

# We call our custom layer on the input and the decoded output,
# to obtain the final model output.
y = CustomVariationalLayer()([input_img, z_decoded])
*****
vae = Model(input_img, y)
vae.compile(optimizer='rmsprop', loss=None)
vae.summary()

# Train the VAE on MNIST digits
(x_train, _), (x_test, y_test) = mnist.load_data()

x_train = x_train.astype('float32') / 255.
x_train = x_train.reshape(x_train.shape + (1,))
x_test = x_test.astype('float32') / 255.
x_test = x_test.reshape(x_test.shape + (1,))

vae.fit(x=x_train, y=None,
        shuffle=True,
        epochs=10,
        batch_size=batch_size,
        validation_data=(x_test, None))

*****
import matplotlib.pyplot as plt
from scipy.stats import norm

```

```

# Display a 2D manifold of the digits
n = 15 # figure with 15x15 digits
digit_size = 28
figure = np.zeros((digit_size * n, digit_size * n))
# Linearly spaced coordinates on the unit square were transformed
# through the inverse CDF (ppf) of the Gaussian
# to produce values of the latent variables z,
# since the prior of the latent space is Gaussian
grid_x = norm.ppf(np.linspace(0.05, 0.95, n))
grid_y = norm.ppf(np.linspace(0.05, 0.95, n))

for i, yi in enumerate(grid_x):
    for j, xi in enumerate(grid_y):
        z_sample = np.array([[xi, yi]])
        z_sample = np.tile(z_sample, batch_size).reshape(batch_size, 2)
        x_decoded = decoder.predict(z_sample, batch_size=batch_size)
        digit = x_decoded[0].reshape(digit_size, digit_size)
        figure[i * digit_size: (i + 1) * digit_size,
              j * digit_size: (j + 1) * digit_size] = digit

plt.figure(figsize=(10, 10))
plt.imshow(figure, cmap='Greys_r')
plt.show()
'''

```

[4]: "\nExample Code from 8.4 in Deep Learning with Python provided below for reference:\n*****\n\nEncode the input into a mean and variance parameter\nnz_mean, z_log_variance = encoder(input_img)\n\n# Draw a latent point using a small random epsilon\nnz = z_mean + exp(z_log_variance) * epsilon\n\n# Then decode z back to an image\nreconstructed_img = decoder(z)\n\n# Instantiate a model\nmodel = Model(input_img, reconstructed_img)\n\n# Then train the model using 2 losses:\n# a reconstruction loss and a regularization loss\n*****\nimg_shape = (28, 28, 1)\nbatch_size = 16\nlatent_dim = 2 # Dimensionality of the latent space: a plane\n\ninput_img = keras.Input(shape=img_shape)\n\nx = layers.Conv2D(32, 3,\n padding='same',\n activation='relu')(input_img)\n\nx = layers.Conv2D(64, 3,\n padding='same', activation='relu',\n strides=(2, 2))(x)\n\nx = layers.Conv2D(64, 3,\n padding='same', activation='relu')(x)\n\nx_shape_before_flattening = K.int_shape(x)\n\nx = layers.Flatten()(x)\n\nx = layers.Dense(32, activation='relu')(x)\n\nnz_mean = layers.Dense(latent_dim)(x)\n\nnz_log_var = layers.Dense(latent_dim)(x)\n*****\n*****\n\ndef sampling(args):\n z_mean, z_log_var = args\n epsilon =


```

decoder.predict(z_sample, batch_size=batch_size)\n            digit =
x_decoded[0].reshape(digit_size, digit_size)\n            figure[i * digit_size: (i
+ 1) * digit_size,\n                    j * digit_size: (j + 1) * digit_size] =
digit\n\nplt.figure(figsize=(10, 10))\nplt.imshow(figure,
cmap='Greys_r')\nplt.show()\n"

```

0.2.1 Create Directories

```

[5]: results_dir = Path('C:/Users/jkmey/Documents/Github/DSC650_Course_Assignments/
    ↪dsc650/dsc650/assignments/assignment12/').joinpath('results').joinpath('vae')
results_dir.mkdir(parents = True, exist_ok = True)

```

0.2.2 Create Encoder Network (Simple Convnet)

```

[6]: '''
Encoder network: A simple convnet which maps the input image x to two vectors,
    ↪z_mean and z_log_variance.
'''

img_shape = (28, 28, 1)
batch_size = 16
latent_dim = 2 # Dimensionality of the latent space: a plane

input_img = keras.Input(shape=img_shape)

x = layers.Conv2D(32, 3,
                  padding='same', activation='relu')(input_img)
x = layers.Conv2D(64, 3,
                  padding='same', activation='relu',
                  strides=(2, 2))(x)
x = layers.Conv2D(64, 3,
                  padding='same', activation='relu')(x)
x = layers.Conv2D(64, 3,
                  padding='same', activation='relu')(x)
shape_before_flattening = K.int_shape(x)

x = layers.Flatten()(x)
x = layers.Dense(32, activation='relu')(x)

z_mean = layers.Dense(latent_dim)(x)
z_log_var = layers.Dense(latent_dim)(x)

```

0.2.3 Create Sampling Function

```

[7]: '''
Here is the code for using z_mean and z_log_var, the parameters of the
    ↪statistical distribution assumed to

```

have produced `input_img`, to generate a latent space point `z`. Here, we wrap some
↳ arbitrary code
(built on top of Keras backend primitives) into a `Lambda` layer. In Keras,
↳ everything needs to be a layer,
so code that isn't part of a built-in layer should be wrapped in a `Lambda` (or
↳ else, in a custom layer).

```
'''
def sampling(args):
    z_mean, z_log_var = args
    epsilon = K.random_normal(shape=(K.shape(z_mean)[0], latent_dim),
                              mean=0., stddev=1.)
    return z_mean + K.exp(z_log_var) * epsilon

z = layers.Lambda(sampling)([z_mean, z_log_var])
```

0.2.4 Decoder Implementation

```
[8]: '''
This is the decoder implementation: we reshape the vector z to the dimensions
↳ of an image,
then we use a few convolution layers to obtain a final image output that has
↳ the same dimensions
as the original input_img.
'''

# This is the input where we will feed `z`.
decoder_input = layers.Input(K.int_shape(z)[1:])

# Upsample to the correct number of units
x = layers.Dense(np.prod(shape_before_flattening[1:]),
                  activation='relu')(decoder_input)

# Reshape into an image of the same shape as before our last `Flatten` layer
x = layers.Reshape(shape_before_flattening[1:])(x)

# We then apply then reverse operation to the initial
# stack of convolution layers: a `Conv2DTranspose` layers
# with corresponding parameters.
x = layers.Conv2DTranspose(32, 3,
                           padding='same', activation='relu',
                           strides=(2, 2))(x)

x = layers.Conv2D(1, 3,
                  padding='same', activation='sigmoid')(x)

# We end up with a feature map of the same size as the original input.

# This is our decoder model.
decoder = Model(decoder_input, x)
```



```
# We then apply it to `z` to recover the decoded `z`.
z_decoded = decoder(z)
```

0.2.5 Create Custom Variational Layer Class

```
[9]: '''
The dual loss of a VAE doesn't fit the traditional expectation of a sample-wise
    ↪ function of the form loss(input, target).
Thus, we set up the loss by writing a custom layer with internally leverages
    ↪ the built-in add_loss layer method to create
an arbitrary loss.
'''
class CustomVariationalLayer(keras.layers.Layer):

    def vae_loss(self, x, z_decoded):
        x = K.flatten(x)
        z_decoded = K.flatten(z_decoded)
        xent_loss = keras.metrics.binary_crossentropy(x, z_decoded)
        kl_loss = -5e-4 * K.mean(
            1 + z_log_var - K.square(z_mean) - K.exp(z_log_var), axis=-1)
        return K.mean(xent_loss + kl_loss)

    def call(self, inputs):
        x = inputs[0]
        z_decoded = inputs[1]
        loss = self.vae_loss(x, z_decoded)
        self.add_loss(loss, inputs=inputs)
        # We don't use this output.
        return x

# We call our custom layer on the input and the decoded output,
# to obtain the final model output.
y = CustomVariationalLayer()([input_img, z_decoded])
```

0.2.6 Instantiate and Train the Model

```
[10]: '''
Finally, we instantiate and train the model. Since the loss has been taken care
    ↪ of in our custom layer,
we don't specify an external loss at compile time (loss=None),
which in turns means that we won't pass target data during training
(as you can see we only pass x_train to the model in fit).
'''
vae = Model(input_img, y)
vae.compile(optimizer='rmsprop', loss=None)
vae.summary()
```

```
# Train the VAE on MNIST digits
(x_train, _), (x_test, y_test) = mnist.load_data()

x_train = x_train.astype('float32') / 255.
x_train = x_train.reshape(x_train.shape + (1,))
x_test = x_test.astype('float32') / 255.
x_test = x_test.reshape(x_test.shape + (1,))

vae.fit(x=x_train, y=None,
        shuffle=True,
        epochs=10,
        batch_size=batch_size,
        validation_data=(x_test, None))
```

WARNING:tensorflow:Output {0} missing from loss dictionary. We assume this was done on purpose. The fit and evaluate APIs will not be expecting any data to be passed to custom_variational_layer.

Model: "model_1"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 28, 28, 1)]	0	[]
conv2d (Conv2D) ['input_1[0][0]']	(None, 28, 28, 32)	320	
conv2d_1 (Conv2D) ['conv2d[0][0]']	(None, 14, 14, 64)	18496	
conv2d_2 (Conv2D) ['conv2d_1[0][0]']	(None, 14, 14, 64)	36928	
conv2d_3 (Conv2D) ['conv2d_2[0][0]']	(None, 14, 14, 64)	36928	
flatten (Flatten) ['conv2d_3[0][0]']	(None, 12544)	0	
dense (Dense) ['flatten[0][0]']	(None, 32)	401440	
dense_1 (Dense)	(None, 2)	66	['dense[0][0]']
dense_2 (Dense)	(None, 2)	66	['dense[0][0]']
lambda (Lambda)	(None, 2)	0	

```

['dense_1[0][0]',
'dense_2[0][0]']

model (Functional)                (None, 28, 28, 1)    56385
['lambda[0][0]']

custom_variational_layer (Cust    (None, 28, 28, 1)    0
['input_1[0][0]',
omVariationalLayer)                'model[0][0]']

```

```

=====
=====

```

```

Total params: 550,629
Trainable params: 550,629
Non-trainable params: 0

```

```

-----
-----

```

```

Train on 60000 samples, validate on 10000 samples

```

```

Epoch 1/10

```

```

WARNING:tensorflow:OMP_NUM_THREADS is no longer used by the default Keras
config. To configure the number of threads, use tf.config.threading APIs.
59984/60000 [=====>.] - ETA: 0s - loss: 0.2098

```

```

C:\Users\jkmey\anaconda3\envs\dsc650\lib\site-
packages\keras\engine\training_v1.py:2333: UserWarning: `Model.state_updates`
will be removed in a future version. This property should not be used in
TensorFlow 2.0, as `updates` are applied automatically.

```

```

    updates = self.state_updates

```

```

60000/60000 [=====] - 113s 2ms/sample - loss: 0.2098 -
val_loss: 0.1952

```

```

Epoch 2/10

```

```

60000/60000 [=====] - 112s 2ms/sample - loss: 0.1921 -
val_loss: 0.1893

```

```

Epoch 3/10

```

```

60000/60000 [=====] - 115s 2ms/sample - loss: 0.1880 -
val_loss: 0.1879

```

```

Epoch 4/10

```

```

60000/60000 [=====] - 112s 2ms/sample - loss: 0.1856 -
val_loss: 0.1840

```

```

Epoch 5/10

```

```

60000/60000 [=====] - 117s 2ms/sample - loss: 0.1840 -
val_loss: 0.1836

```

```

Epoch 6/10

```

```

60000/60000 [=====] - 112s 2ms/sample - loss: 0.1828 -
val_loss: 0.1822

```

```

Epoch 7/10

```

```

60000/60000 [=====] - 112s 2ms/sample - loss: 0.1818 -
val_loss: 0.1836

```

```
Epoch 8/10
60000/60000 [=====] - 114s 2ms/sample - loss: 0.1811 -
val_loss: 0.1822
Epoch 9/10
60000/60000 [=====] - 116s 2ms/sample - loss: 0.1804 -
val_loss: 0.1809
Epoch 10/10
60000/60000 [=====] - 112s 2ms/sample - loss: 0.1799 -
val_loss: 0.1800
```

```
[10]: <keras.callbacks.History at 0x260465059f0>
```

0.2.7 Generate Plot and Save as Output to Results Directory

```
[11]: '''
Once such a model is trained -- e.g. on MNIST, in our case --
we can use the decoder network to turn arbitrary latent space vectors into
→images:
'''
# Display a 2D manifold of the digits
n = 15 # figure with 15x15 digits
digit_size = 28
figure = np.zeros((digit_size * n, digit_size * n))
# Linearly spaced coordinates on the unit square were transformed
# through the inverse CDF (ppf) of the Gaussian
# to produce values of the latent variables z,
# since the prior of the latent space is Gaussian
grid_x = norm.ppf(np.linspace(0.05, 0.95, n))
grid_y = norm.ppf(np.linspace(0.05, 0.95, n))

for i, yi in enumerate(grid_x):
    for j, xi in enumerate(grid_y):
        z_sample = np.array([[xi, yi]])
        z_sample = np.tile(z_sample, batch_size).reshape(batch_size, 2)
        x_decoded = decoder.predict(z_sample, batch_size=batch_size)
        digit = x_decoded[0].reshape(digit_size, digit_size)
        figure[i * digit_size: (i + 1) * digit_size,
              j * digit_size: (j + 1) * digit_size] = digit

plt.figure(figsize=(10, 10))
plt.imshow(figure, cmap='Greys_r')
digit_image = results_dir.joinpath('assignment12_MeyerJake_15x15_Digit_Grid.
→png')
plt.savefig(digit_image)
plt.show()
```

```
C:\Users\jkmey\anaconda3\envs\dsc650\lib\site-
packages\keras\engine\training_v1.py:2357: UserWarning: `Model.state_updates`
```

will be removed in a future version. This property should not be used in TensorFlow 2.0, as `updates` are applied automatically.

```
updates=self.state_updates,
```

